

# Enterprise Customer Onboarding Agent — Solution Design Document

February 2025 — StackAdapt Case Study Submission

## 1. Architecture Overview

The solution implements an **AI-powered automation agent** that orchestrates enterprise customer onboarding from deal closure through SaaS provisioning. Built with **LangGraph** for state machine orchestration and **FastAPI** for the REST interface, the agent integrates with multiple enterprise systems, validates business rules, assesses risks using LLM intelligence, and takes autonomous actions including tenant provisioning and task management.

### 1.1 System Integrations

The agent connects to enterprise systems via REST APIs. Field selection focused on business-critical data to minimize payload and optimize validation:

| System         | Objects & Key Fields   |
|----------------|--|
| Salesforce CRM | Account (Id, Name, IsDeleted, Status, BillingCountry), Opportunity (StageName, Amount, CloseDate), User (IsActive, Email). <a href="#">[API Ref]</a> |
| NetSuite ERP   | Invoice (status, dueDate, amountRemaining, paymentStatus). Used for payment verification before provisioning. <a href="#">[API Ref]</a>              |
| CLM System     | Contract status, signatories, effective dates, key terms. Mock implementation simulating DocuSign CLM.   |
| Provisioning   | Internal system for tenant creation + 14-task onboarding checklist with dependencies and due dates.  |

**Field Selection Rationale:** `Account.IsDeleted` validates account exists; `Opportunity.StageName="Closed Won"` confirms deal closure; `Invoice.status` identifies payment blockers. This minimizes API calls while capturing decision-critical data.




Figure 1: High-level architecture showing trigger sources, LangGraph orchestration, and integration layer.

## 2. AI Agent Application

The agent leverages **OpenAI GPT-4** for intelligent analysis with a deterministic rule-based fallback, ensuring the system operates even without LLM connectivity. This dual-mode approach balances intelligence with reliability.

### 2.1 LLM-Powered Intelligence

- Risk Analysis:** Evaluates API errors, business rule violations, and warnings to generate risk levels (low/medium/high/critical) with business impact assessment and estimated resolution time.
- Summary Generation:** Creates human-readable status reports for Customer Success teams, translating technical states into actionable insights.
- Action Recommendations:** Produces prioritized remediation steps with owner assignment (CS, Finance, IT, Legal) based on issue type and urgency.
- Error Interpretation:** Converts technical API error codes (e.g., `INVALID_SESSION_ID`, `INSUFFICIENT_ACCESS`) into plain-English explanations with specific resolution steps.

### 2.2 Autonomous Actions

Upon decision, the agent executes: **(1)** Auto-provisions tenant with tier-appropriate configuration; **(2)** Creates 14-task onboarding checklist with dependencies, owners, and due dates; **(3)** Sends Slack alerts to `#cs-onboarding` channel; **(4)** Emails welcome message with login credentials to customer; **(5)** Records all API failures with full context for audit.

## 3. Orchestration & Event-Driven Flows

The agent supports multiple trigger mechanisms for different operational needs. LangGraph manages state transitions through a defined workflow with full observability via structured logging and optional LangSmith tracing.




Figure 2: Decision logic: API Errors → BLOCK, Violations → BLOCK, Warnings → ESCALATE, Clear → PROCEED.

| Trigger     | Source          | Endpoint                    |
|-------------|-----------------|-----------------------------|
| Webhook     | Salesforce Flow | POST /webhook/onboarding    |
| Manual      | CS Team         | POST /demo/run/{account.id} |
| Batch       | Scheduler/Cron  | POST /demo/run-all          |
| Task Update | CS Action       | PUT /demo/tasks/{id}/{task} |

**State Machine Flow:** Initialize → Fetch Salesforce → Fetch CLM → Fetch NetSuite → Validate Invariants → Analyze Risks (LLM) → Make Decision → [Provision if PROCEED] → Send Notifications → Generate Summary → Complete. Each node handles errors gracefully and updates state.

## 4. Trade-offs, Assumptions & Considerations

### 4.1 Design Decisions

| Decision               | Rationale   |
|------------------------|---|
| API Errors → BLOCK     | Data integrity over speed; cannot provision with incomplete data        |
| Synchronous Processing | Simplicity for demo; production would use message queues (SQS/RabbitMQ) |
| Rule-based Fallback    | Ensures availability without LLM; deterministic behavior for testing    |
| 14 Fixed Tasks         | Predictable workflow; production would use configurable templates       |

## 5. Multi-Agent Collaboration via Model Context Protocol (MCP)

The architecture supports multi-agent collaboration through MCP, enabling specialized agents to communicate via standardized tool interfaces. Each agent is a domain expert with focused responsibilities.

| Agent          | Responsibility & MCP Tools   |
|----------------|--|
| Coordinator    | Orchestrates end-to-end workflow ( <code>salesforce.*</code> , <code>provision.*</code> , <code>tasks.*</code> ) |
| Contract Agent | Monitors signature status, sends reminders ( <code>clm.get_contract</code> , <code>clm.send_reminder</code> )    |
| Finance Agent  | Tracks payments, initiates dunning ( <code>netsuite.get_invoice</code> , <code>netsuite.send_dunning</code> )    |
| Task Monitor   | Detects overdue tasks, alerts CS ( <code>tasks.get_overdue</code> , <code>notify.alert</code> )                  |

**Benefits:** Separation of concerns (domain experts) • Reusable integrations (shared MCP servers) • Scalable complexity (add agents without modifying existing) • Cross-agent context sharing via tool responses.

**Error Simulation for Resilience Testing:** The endpoint `/demo/enable-random-errors` injects configurable failures with adjustable rates:


- Authentication errors (HTTP 401) - expired tokens
- Validation errors (HTTP 400) - malformed data
- Rate limit errors (HTTP 429) - API throttling
- Server errors (HTTP 500) - backend failures

This enables chaos testing without modifying business logic, validating the agent's error handling and recovery capabilities.

### 4.2 Scalability & Security

**Scalability Path:** Current demo uses synchronous processing with in-memory state. Production deployment would include: message queues for async processing, Redis for distributed state, Kubernetes for horizontal scaling, and LLM response caching for cost optimization.

**Security & Governance:** OAuth 2.0 with token refresh simulation, structured audit trails with correlation IDs across all operations, PII masking in logs, role-based access control (RBAC) for API operations, and complete state snapshots for compliance.



**Figure 3:** MCP architecture: Specialized agents collaborate through shared MCP servers wrapping external APIs.