Enterprise Customer Onboarding Agent — Solution Design Document

April 2026

1. Architecture Overview

The solution implements an **AI-powered automation agent** that orchestrates enterprise customer onboarding from deal closure through SaaS provisioning. Built with **Pydantic AI** for agentic tool calling and **FastAPI** for the REST interface, the agent uses native LLM tool calling to autonomously reason through 24 registered tools. Each tool is mirrored as a **FastMCP** server definition for MCP-style extraction. A second **CS Assistant Agent** shares all tools but returns free-form text for conversational interaction via a **Streamlit UI** with four pages (Dashboard, Portfolio Overview, Run Onboarding, Chat). Data flows sequentially across systems via **cross-system chaining**: Opportunity.ContractId → SF Contract.Id → CLM.salesforce_contract_id → Invoice.clmContractRef, with fallback to account-based lookups when chain links are missing.

1.1 System Integrations

System	Objects & Key Fields
Salesforce CRM	`Account` (Id, Name, IsDeleted, OwnerId), `Opportunity` (StageName, Amount, **ContractId**), `Contract` (**Id**, Status, StartDate), `User` (IsActive, Email). [API Ref]
CLM System	contract_id, status, **salesforce_contract_id** (chains from SF Contract.Id), signatories, key_terms. Mock simulating DocuSign CLM.
NetSuite ERP	`Invoice` (status, dueDate, amountRemaining, total, **clmContractRef** chains from CLM contract_id). [API Ref]
Provisioning	Tenant creation + 14-task onboarding checklist with dependencies, owners, and due dates.
Currency (Frankfurter)	Live and historical exchange rates via ECB data. Converts invoice amounts across currencies (e.g., CAD↔USD) using the invoice date rate. [API Ref]

Field Selection Rationale: `Account.IsDeleted` validates account exists; `Opportunity.StageName="Closed Won"` confirms deal closure; `Invoice.status` identifies payment blockers; `CLM.status` in {EXECUTED, SIGNED} required for provisioning. Currency conversion uses historical rates from the invoice date for financial accuracy. Minimizes API calls while capturing decision-critical data.

1.2 Data Flow

(1) Webhook with `account_id` + `correlation_id` **(2)** Fetch account, then user/opportunity in parallel; chain sequentially: Opp.ContractId → SF Contract → CLM (via salesforce_contract_id) → Invoice (via clmContractRef) **(3)** `validate_business_rules` + `check_financial_alignment` **(4)** PROCEED → provision + monitor; BLOCK/ESCALATE → notify.

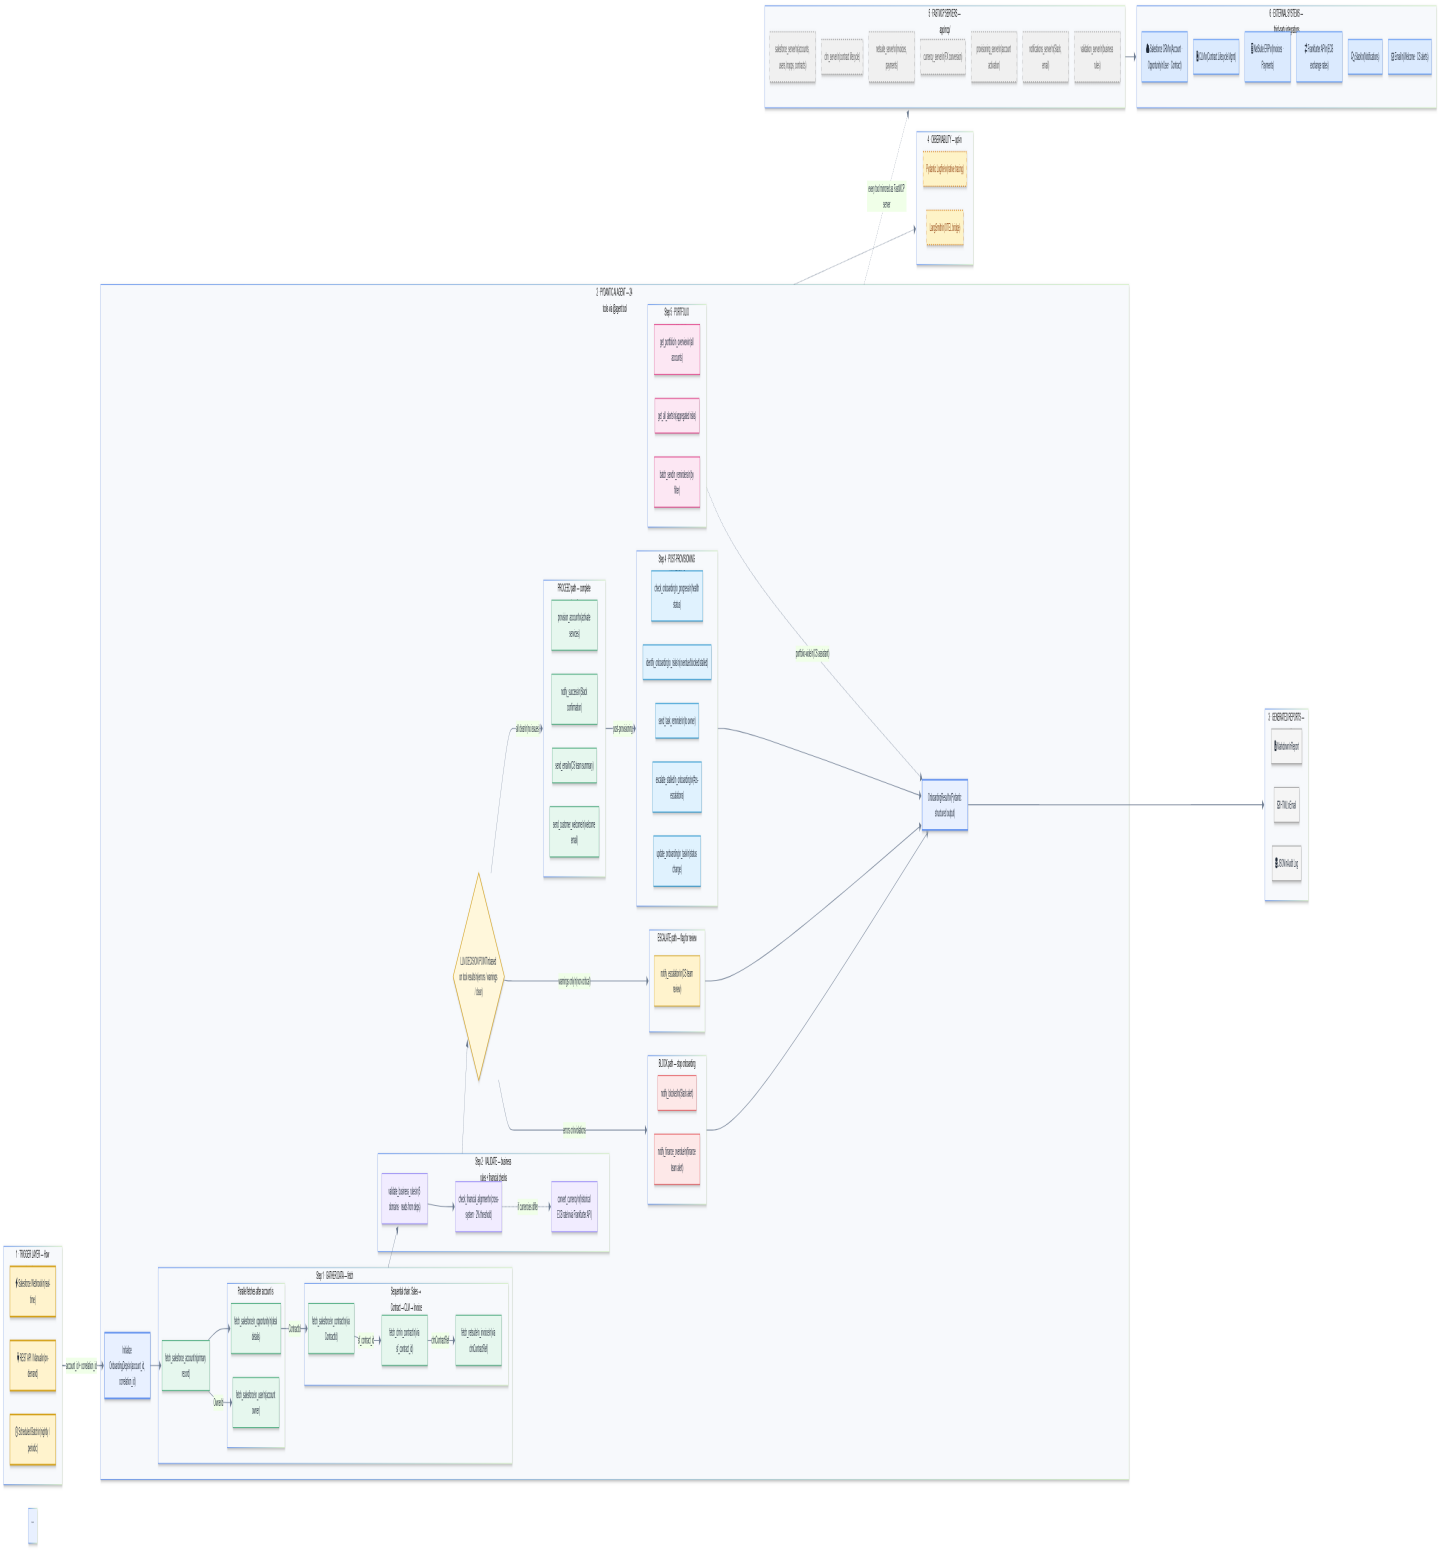

Figure 1: Architecture: triggers, Pydantic AI agent orchestration with 24 tools, cross-system data chaining, FastMCP servers, and external system integrations.

2. AI Agent Application

The agent leverages **OpenAI GPT-4o-mini** (configurable via `OPENAI_MODEL`) with **Pydantic AI** for native tool calling. The LLM autonomously decides which of the 24 registered tools to call, in what order, and how many times. Falls back to **Ollama** local models when no API key is available. A second **CS Assistant Agent** shares all tools with free-form text output for conversational monitoring via the Streamlit UI.

2.1 LLM-Powered Intelligence

- **Risk Analysis**: Evaluates API errors, violations, warnings → risk levels (low/medium/high/critical) with business impact and resolution time estimates.
- **Summary Generation**: Human-readable reports for CS teams, translating technical states to actionable insights.
- **Action Recommendations**: Prioritized steps with owner assignment (CS, Finance, IT/DevOps, Legal, Sales Ops).
- **Error Interpretation**: Converts technical codes (e.g., `INVALID_SESSION_ID`) to plain English with resolution steps.

2.2 Autonomous Actions

(1) Auto-provisions the SaaS tenant with tier-appropriate configuration (Enterprise/Growth/Starter).
(2) Creates a 14-task onboarding checklist with dependencies, ownership (system/cs_team/customer), and due dates.
(3) Sends Slack notifications to `#cs-onboarding-alerts` (blocked) or `#cs-onboarding` (success/escalate).
(4) Sends customer-facing welcome emails with tenant ID and login URL.
(5) Generates HTML/Markdown/JSON reports for auditing and traceability.
(6) Post-provisioning: monitors progress, detects risks, sends reminders, escalates stalled onboardings.
(7) Portfolio management: aggregated alerts across accounts, suggested actions with approve/dismiss, batch reminders.

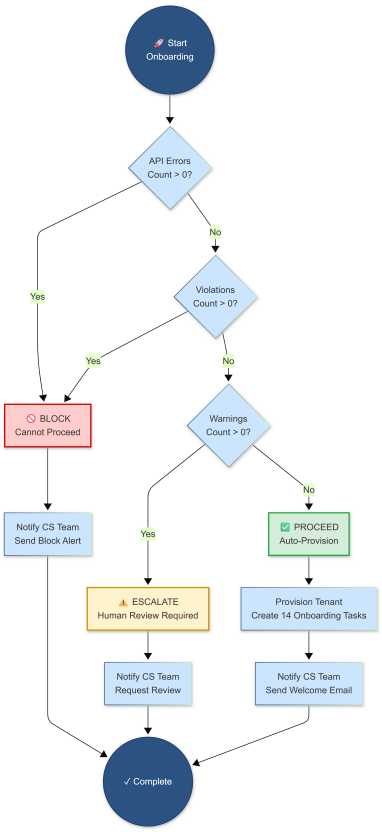

Figure 2: Decision: API Errors/Violations → BLOCK, Warnings → ESCALATE, Clear → PROCEED.

2.3 Onboarding Task Management

The provisioning system creates a 14-task onboarding checklist:

Task Categories:

- **Automated (4)**: Tenant creation, API credentials, welcome email, training materials
- **CS Team (5)**: Schedule/conduct kickoff, SSO config, custom reports, 30-day check-in
- **Customer (4)**: Verify login, complete tour, invite team, create first campaign
- **Milestone (1)**: Onboarding complete marker

API Endpoints:

- `GET /tasks/{id}` – All tasks with status
- `GET /tasks/{id}/pending?owner=cs_team` – Filtered pending
- `GET /tasks/{id}/overdue` – Overdue task alerts
- `PUT /tasks/{id}/{task_id}` – Update task status
- `GET /tasks/{id}/next-actions` – Ready-to-execute tasks

Tasks include dependencies (e.g., "Complete Tour" depends on "Verify Login"), due dates calculated relative to provisioning, and ownership assignment enabling filtered views per team.

2.4 Business Rule Validation (Invariants)

The agent enforces a two-tier validation system across five domains. **Violations** are blocking (cause BLOCK decision); **Warnings** are non-blocking (cause ESCALATE for human review).

Domain	Blocking Violations	Non-blocking Warnings
Account	Missing Id/Name, IsDeleted=true	Missing BillingCountry, Industry, OwnerId
Opportunity	Stage ≠ "Closed Won", Missing Id/AccountId	Missing Amount, CloseDate, ContractId
Contract (CLM)	Status ∉ {EXECUTED, SIGNED}, Missing contract_id	No effective/expiry date, pending signatories
Invoice	Voided, Cancelled, Missing invoice_id	Overdue, Pending, <50% paid, missing due_date
User	Inactive, Missing Id/Email/ProfileId	Missing Title, Department, ManagerId

3. Orchestration & Event-Driven Flows

Trigger	Source	Endpoint
Webhook	Salesforce Flow	`POST /webhook/onboarding`
Manual	CS Team	`POST /demo/run/{account_id}`
Batch	Scheduler	`POST /demo/run-all`

Error Simulation: `/demo/enable-random-errors` injects failures (401 auth, 400 validation, 429 rate limit, 500 server) at configurable rates for chaos testing without modifying business logic.

Report Generation: Each run produces three output files:

- `run_report_*.md` – Full Markdown report with all details
- `email_*.html` – Professional HTML email (blocked, escalation, success, or welcome)
- `audit_*.json` – Machine-readable audit log

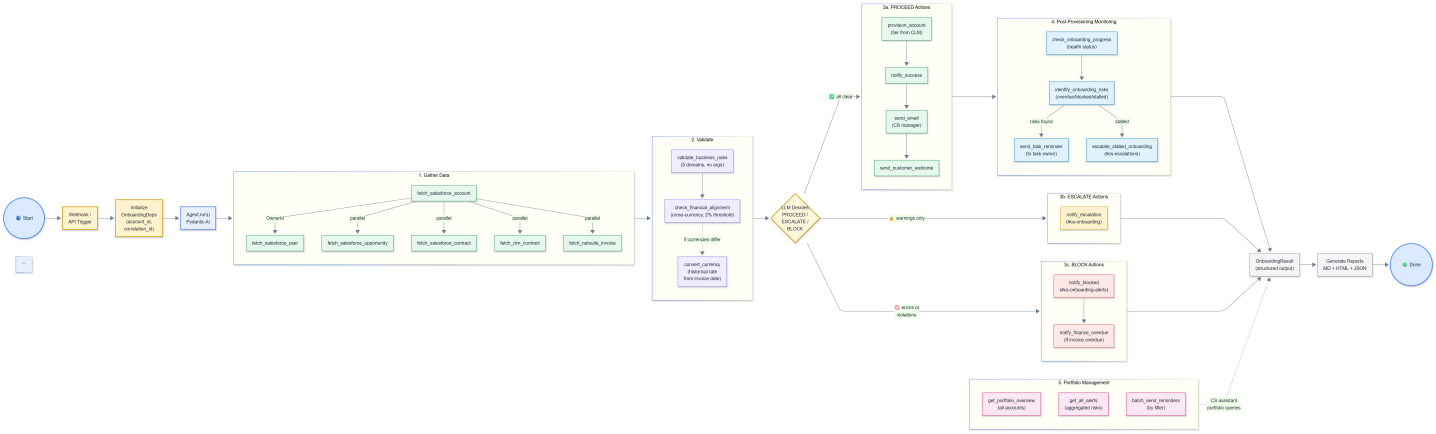

Figure 3: Agent workflow: full lifecycle from webhook ingestion through tool calling, validation, decision routing, post-provisioning monitoring, and portfolio management.

3.1 API Error Classification & Handling

Each integration returns structured error payloads with `error_type`, `error_code`, `http_status`, `message`, `resolution`, and `owner`:

Error Type	HTTP	Resolution
Authentication	401	Re-authenticate, refresh token (IT/DevOps)
Authorization	403	Check profile/permission sets (System Admin)
Validation	400	Fix field values in source (Data Admin)
Rate Limit	429	Implement backoff, increase quota (Integration Admin)
Server Error	500	Check status page, retry later (Support Team)

3.2 Error Handling Features

- **In-place Error Simulator**: `ERROR_SIMULATOR` object modified in-place when enabled; all modules reference the same instance
- **Comprehensive Catching**: Each integration catches specific error types AND generic `APIError` as fallback
- **Error-Aware Decisions**: `make_decision` checks `api_errors` first — system failures block onboarding
- **Full Traceability**: API errors added to violations and appear in reports with context

4. Trade-offs, Assumptions & Considerations

4.1 Key Trade-offs

Decision	Trade-off
API Errors → BLOCK	Integrity over speed
Sequential Data Fetching	Clearer failure tracing but has higher latency
Rule-based Fallback	Less smart but reliable
14 Fixed Tasks	Predictable but rigid
LLM Advisory Only	Auditable but less adaptive

4.2 Assumptions

- Salesforce is source of truth for accounts
- CLM status reflects actual signatures
- Invoice payment status is current
- Cross-currency comparison uses historical ECB rates from invoice date with 2% threshold
- One invoice per account per onboarding run
- One Opportunity per Account (no multi-deal)
- Contract has a single version, no amendments
- CS monitors Slack channels for alerts
- Single onboarding per account at a time
- Customer email available in CLM signatories
- Account ID consistent across all systems

4.3 Limitations

- **Mock integrations**: Salesforce, NetSuite, CLM use simulated data; not connected to real APIs
- **In-memory state**: Provisioning and task data lost on restart; no persistent database
- **Single instance**: No horizontal scaling or load balancing
- **Limited approval UI**: Suggested Actions panel provides approve/dismiss; full Slack interactive approval not yet implemented
- **No webhook auth**: Inbound webhooks not verified via HMAC or shared secret

4.4 Testing & Demonstration

The system includes 10 pre-configured demo scenarios validating all decision paths. Three PROCEED scenarios include post-provisioning simulation to demonstrate proactive risk detection:

Account ID	Scenario	Description	Decision
`ACME-001`	Happy Path	All checks pass, contract executed, invoice paid	PROCEED
`BETA-002`	Opportunity Blocked	Opportunity in "Negotiation" stage	BLOCK
`GAMMA-003`	Overdue Invoice	Invoice 25+ days overdue, needs finance review	ESCALATE
`DELETED-004`	Deleted Account	Account marked `IsDeleted=true`	BLOCK
`MISSING-999`	Account Not Found	Account does not exist in any system	BLOCK
`FOREX-005`	FX Invoice Mismatch	CAD $145k invoice vs USD $100k opportunity; historical rate gap >2%	ESCALATE
`PARTIAL-006`	Partial Payment Gap	$190k paid of $200k invoice (5% underpayment exceeds 2% threshold)	ESCALATE
`STARTER-007`	Customer Not Logged In	All checks pass; post-provisioning: customer hasn't logged in after 5 days	PROCEED
`GROWTH-008`	Stalled Onboarding	All checks pass; post-provisioning: <30% complete after 10 days, kickoff not done	PROCEED
`ENTERPRISE-009`	SSO & Blocked Tasks	All checks pass; post-provisioning: SSO blocked, multiple tasks overdue after 8 days	PROCEED

4.5 Observability & Tracing

The agent supports **dual opt-in tracing** — both can run simultaneously:

- **Pydantic Logfire**: Native agent tracing — reasoning chains, tool calls with inputs/outputs, structured output validation.
- **LangSmith**: Agent execution runs, tool invocations, LLM completions via OpenTelemetry bridge (`OtelSpanProcessor`). Logfire owns the TracerProvider; LangSmith attaches as a span processor.
- **Structured JSON Logging**: All events logged with `correlation_id`, timestamps, and context.

4.6 Security Patterns

All integrations are mocked but demonstrate production security patterns:

- **OAuth Simulation**: Token expiry and refresh flows with distinct error types for expired vs invalid tokens
- **Permission Checking**: Role-based access control verified before every API call
- **Credential Validation**: Structured error hierarchy (`AuthenticationError`, `AuthorizationError`, `ValidationError`, `RateLimitError`)
- **Audit Logging**: Correlation IDs propagated end-to-end for traceability across all systems
- **Error Masking**: No sensitive data (tokens, secrets) exposed in API responses or reports

5. Multi-Agent Architecture via Model Context Protocol (MCP) Server Connections

Note: All 24 agent tools are mirrored as FastMCP server definitions in `app/mcp/` (7 servers: Salesforce, CLM, NetSuite, Currency, Provisioning, Notifications, Validation). These can be extracted to standalone MCP services for multi-agent collaboration.

The architecture supports multi-agent collaboration through the **Model Context Protocol (MCP)**, enabling specialized agents to interact through standardized, governed tool interfaces rather than direct API access. Each integration is wrapped as a FastMCP server, cleanly separating decision logic from execution concerns.

Agent	Responsibility & MCP Tools
Coordinator	Orchestrates end-to-end workflow (`salesforce.*`, `provision.*`)
Contract Agent	Monitors contract execution and signatures (`clm.get_contract`)
Finance Agent	Evaluates billing, payments, and cross-currency alignment (`netsuite.*`, `currency.*`)
Task Monitor	Detects overdue onboarding tasks (`tasks.get_overdue`)

Benefits: Independent versioning per integration • Clear domain ownership • Centralized access control and audit logging • Scalable agent complexity • Cross-agent context sharing.

Concurrency & API Execution Strategy: MCP servers act as controlled execution layers for external APIs, enabling concurrent and batch-aware request handling. Batch or composite API calls are leveraged where supported (e.g., Salesforce composite reads), with fallback to bounded parallel single-entity requests. Concurrency limits, rate limiting, and circuit breakers are enforced at the MCP layer to prevent downstream overload and cascading failures.

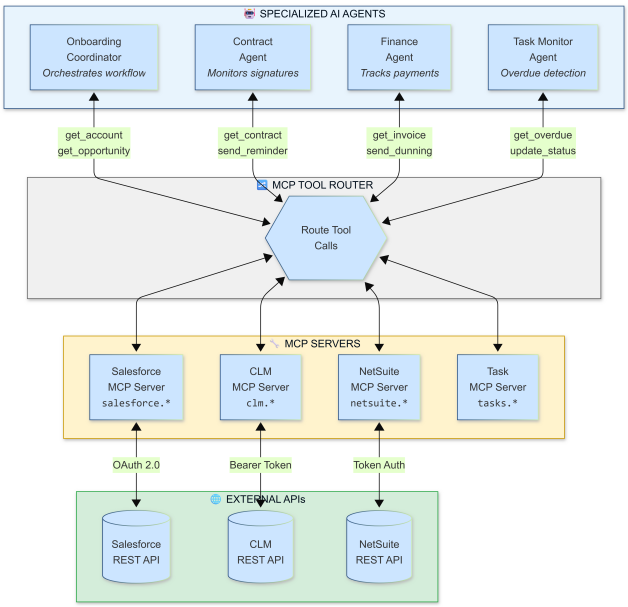

Figure 4: MCP: Specialized agents collaborate through shared execution servers wrapping external APIs.

6. Production Considerations

For additional production enhancements, see the `production_roadmap` directory in the repository.

Cloud Deployment

Frontend: A Streamlit UI is implemented with four pages: Dashboard (proactive alerts, suggested actions with approve/dismiss, results by decision), Portfolio Overview (health distribution, priority actions, sortable account table), Run Onboarding (scenario runner with live provisioning progress), and Chat (portfolio-level commands, account dropdown, batch actions). For production, migrate to a role-based interface (React, ASP.NET) with auth and WebSocket updates.

AWS Infrastructure: Containerized services on ECS/Fargate with horizontal scaling. LLM inference via **AWS Bedrock** for multi-model flexibility (Claude, Llama, Titan), private VPC networking, and IAM-based auth. Bedrock enables cost controls and model switching without code changes.

Docker + Kubernetes: Consistent environments across dev/staging/prod. Auto-scaling for traffic spikes, self-healing services, rolling updates, and zero-downtime deployments.

CI/CD & Observability

CI/CD (GitHub Actions / CodePipeline): Automated pipelines validate agent logic on every PR — unit tests for invariant checks, risk analysis, and decision routing. Deployments are reproducible and auditable with fast feedback loops and safe rollback capability.

Dual Tracing (Implemented): Opt-in dual observability: **Pydantic Logfire** for native agent tracing (reasoning chains, tool calls, structured output validation) and **LangSmith** via OpenTelemetry bridge for LLM execution runs, token usage, and latency. Both run simultaneously — Logfire owns the TracerProvider, LangSmith attaches as a span processor.

Monitoring & Alerting: Prometheus metrics capture throughput, error rates, and decision distributions. Distributed tracing via DataDog or Jaeger for end-to-end visibility. PagerDuty alerts for critical failures.

Agentic Execution & API Governance

Deterministic Decision Boundaries: The LLM reasons through tools via native tool calling (Pydantic AI `@agent.tool` decorators), but deterministic business rules enforced by `validate_business_rules` and `check_financial_alignment` define all violations and warnings. The LLM routes (`PROCEED`/`ESCALATE`/`BLOCK`) based on these outputs, ensuring predictable, testable, auditable outcomes.

Cross-System Data Chaining: Fetch account first, then user/opportunity in parallel. Sequential chain: Opportunity.ContractId → SF Contract → CLM (via salesforce_contract_id) → Invoice (via clmContractRef). Fallback to account-based lookups if chain links are missing. Production optimization: **(1)** Batch APIs (Salesforce Composite); **(2)** Concurrent fetches; **(3)** Bounded exponential backoff (max 3 retries, 1s/2s/4s).

Idempotency: Each run keyed by `correlation_id` propagated across all systems — enables safe replays, deduplication, and re-entrant execution without duplicate side effects.

Error Classification & Retry: Retryable (5xx, rate limits) → bounded exponential backoff. Non-retryable (validation, auth) → immediate halt with diagnostics. Circuit breakers prevent cascading failures.

Rate Limiting: Per-source rate limits and queue-based ingestion for webhook traffic spikes. Per-integration concurrency limits prevent API overload.

Agent Versioning: Every run records agent version, rule set, and prompt version for auditability, regression analysis, and controlled rollouts (canary deployments).

LLM Data Governance: Field-filtered and sanitized data sent to LLMs. PII masked prior to inference. Credentials never exposed beyond integration boundaries.

Multi-Agent Interoperability: All 24 tools mirrored as FastMCP server definitions (7 servers in `app/mcp/`). Extractable to standalone MCP services with independent versioning, access control, and audit logging. A2A protocol for agent-to-agent collaboration.